


02 Linear Predictor



Roadmap

- Linear predictors
- Loss minimization
- Stochastic gradient descent

Application: spam classification

- Input: x = email message


```
From:  pliang@cs.stanford.edu  
Date:  September 26, 2018  
Subject: CS221 announcement
```

```
Hello students,  
I've attached the answers to homework 1...
```

```
From:  a9k62n@hotmail.com  
Date:  September 26, 2018  
Subject: URGENT
```

```
Dear Sir or maDam:  
my friend left sum of 10m dollars...
```

- Output: $y \in \{spam, non - spam\}$
- Objective: obtain a predictor f




Types of prediction tasks

- Binary classification (e.g., email) (spam/not spam):




- Regression (e.g., location, year) (housing price):



Types of prediction tasks


- Multiclass classification: y is a category



- Ranking: y is a permutation



- Structured prediction: y is an object which is built from parts



Question

- Give an example of a prediction task (e.g., image, face/not face).

Data


- Example: species that y is the ground-truth output for x

$$(x, y)$$

- Training data: list of examples

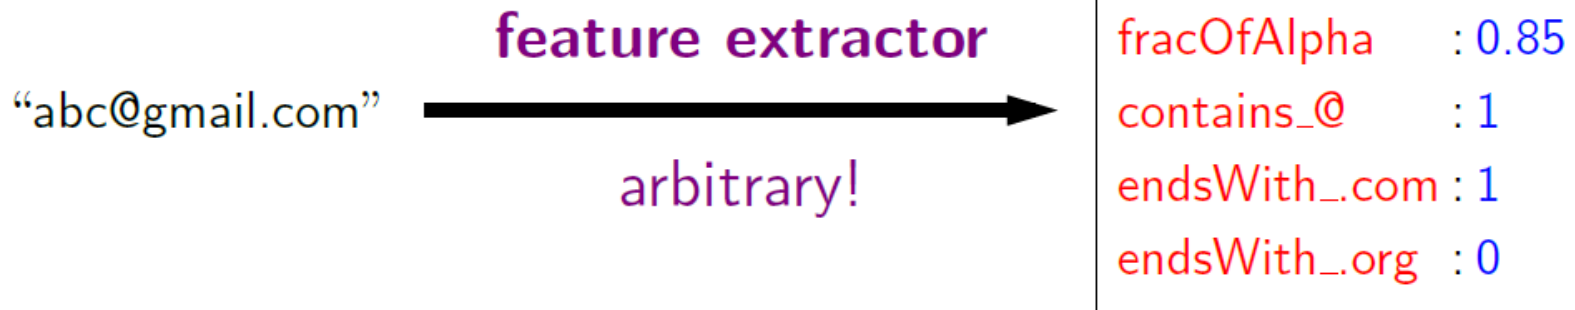
$$\mathcal{D}_{\text{train}} = \begin{bmatrix} (" \dots 10\text{m dollars} \dots", +1), \\ (" \dots \text{CS221} \dots", -1), \\ \end{bmatrix}$$

Framework





Feature extraction

- **Example task:** predict y , whether a string x is an email address
- **Question:** what properties of x might be relevant for predicting y ?
- **Feature extractor:** Given input x , output a set of (feature name, feature value) pairs.



Feature vector notation

- Mathematically, feature vector doesn't need feature names:



Definition: feature vector

For an input x , its feature vector is:

$$\phi(x) = [\phi_1(x), \dots, \phi_d(x)].$$

Think of $\phi(x) \in \mathbb{R}^d$ as a point in a high-dimensional space.

Weight vector

- **Weight vector:** for each feature j , have real number w_j representing contribution of feature to prediction

```
length>10      :-1.2
fracOfAlpha     :0.6
contains_@       :3
endsWith_.com   :2.2
endsWith_.org   :1.4
...
```

Linear predictors

Weight vector $\mathbf{w} \in \mathbb{R}^d$

length>10	:-1.2
fracOfAlpha	:0.6
contains_@	:3
endsWith_.com	:2.2
endsWith_.org	:1.4

Feature vector $\phi(x) \in \mathbb{R}^d$

length>10	:1
fracOfAlpha	:0.85
contains_@	:1
endsWith_.com	:1
endsWith_.org	:0


Score: weighted combination of features

$$\mathbf{w} \cdot \phi(x) = \sum_{j=1}^d w_j \phi(x)_j$$

Example: $-1.2(1) + 0.6(0.85) + 3(1) + 2.2(1) + 1.4(0) = 4.51$

Linear predictors

- Weight vector $\mathbf{w} \in \mathbb{R}^d$
- Feature vector $\phi(x) \in \mathbb{R}^d$
- For binary classification:



Definition: (binary) linear classifier

$$f_{\mathbf{w}}(x) = \text{sign}(\mathbf{w} \cdot \phi(x)) = \begin{cases} +1 & \text{if } \mathbf{w} \cdot \phi(x) > 0 \\ -1 & \text{if } \mathbf{w} \cdot \phi(x) < 0 \\ ? & \text{if } \mathbf{w} \cdot \phi(x) = 0 \end{cases}$$

Geometric intuition

- A binary classifier $f_{\mathbf{w}}$ defines a hyperplane with normal vector \mathbf{w} .

($\mathbb{R}^2 \implies$ hyperplane is a line; $\mathbb{R}^3 \implies$ hyperplane is a plane)

Example:

$$\mathbf{w} = [2, -1]$$



$$\phi(x) \in \{[2, 0], [0, 2], [2, 4]\}$$

[blackboard]


Roadmap

- Linear predictors
- Loss minimization
- Stochastic gradient descent

Framework



Loss functions




Definition: loss function

A loss function $\text{Loss}(x, y, \mathbf{w})$ quantifies how unhappy you would be if you used \mathbf{w} to make a prediction on x when the correct output is y . It is the objective we want to minimize.


Score and margin

- Correct label: y
- Predicted label: $y' = f_{\mathbf{w}}(x) = \text{sign}(\mathbf{w} \cdot \phi(x))$
- Example: $\mathbf{w} = [2, -1], \phi(x) = [2, 0], y = 1$



Definition: score

The score on an example (x, y) is $\mathbf{w} \cdot \phi(x)$, how confident we are in predicting +1



Definition: margin

The margin on an example (x, y) is $\mathbf{w} \cdot \phi(x)y$, how correct we are

Question


- When does a binary classifier predict an error on an example?
 - margin less than 0
 - margin greater than 0
 - score less than 0
 - score greater than 0

Binary classification

Example: $\mathbf{w} = [2, -1]$, $\phi(x) = [2, 0]$, $y = 1$

Recall the binary classifier:

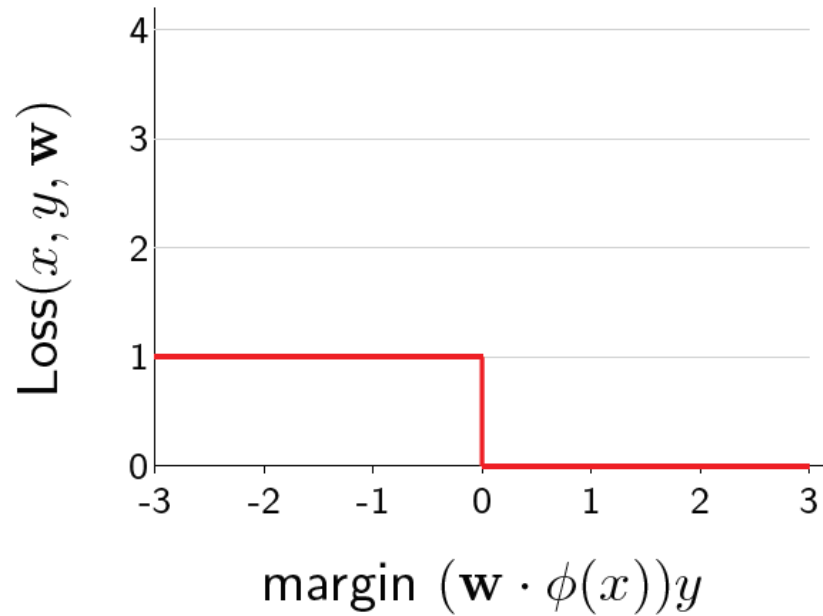
$$f_{\mathbf{w}}(x) = \text{sign}(\mathbf{w} \cdot \phi(x))$$



Definition: zero-one loss

$$\begin{aligned}\text{Loss}_{0-1}(x, y, \mathbf{w}) &= \mathbf{1}[f_{\mathbf{w}}(x) \neq y] \\ &= \mathbf{1}[\underbrace{(\mathbf{w} \cdot \phi(x))y}_{\text{margin}} \leq 0]\end{aligned}$$


Binary classification



$$\text{Loss}_{0-1}(x, y, \mathbf{w}) = \mathbf{1}[(\mathbf{w} \cdot \phi(x))y \leq 0]$$

Linear regression

$$f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x)$$



Definition: residual

The residual is $(\mathbf{w} \cdot \phi(x)) - y$, the amount by which prediction $f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x)$ overshoots the target y .

Linear regression

$$f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x)$$



Definition: squared loss


$$\text{LOSS}_{\text{squared}}(x, y, \mathbf{w}) = \underbrace{(f_{\mathbf{w}}(x) - y)}_{\text{residual}}^2$$

Example :

$$\mathbf{w} = [2, -1], \phi(x) = [2, 0], y = -1$$

$$\text{LOSS}_{\text{squared}}(x, y, \mathbf{w}) = 25$$

Regression loss functions



$$\text{Loss}_{\text{squared}}(x, y, \mathbf{w}) = (\mathbf{w} \cdot \phi(x) - y)^2$$

$$\text{Loss}_{\text{absdev}}(x, y, \mathbf{w}) = |\mathbf{w} \cdot \phi(x) - y|$$

Loss minimization framework

So far: one example, $\text{Loss}(x, y, w)$ is easy to minimize.



Key idea: minimize training loss

$$\text{TrainLoss}(\mathbf{w}) = \frac{1}{|\mathcal{D}_{\text{train}}|} \sum_{(x,y) \in \mathcal{D}_{\text{train}}} \text{Loss}(x, y, \mathbf{w})$$

$$\min_{\mathbf{w} \in \mathbb{R}^d} \text{TrainLoss}(\mathbf{w})$$

Key: need to set w to make global tradeoffs—not every example can be happy.

Which regression loss to use?

Example : $\mathcal{D}_{\text{train}} = \{(1, 0), (1, 2), (1, 1000)\}$

For least squares (L_2) regression :

$$\text{LOSS}_{\text{squared}}(x, y, \mathbf{w}) = (\mathbf{w} \cdot \phi(x) - y)^2$$

- \mathbf{w} that minimizes training loss is **mean** y
- **Mean**: tries to accommodate every example, popular

For least absolute deviation (L_1) regression :

$$\text{LOSS}_{\text{squared}}(x, y, \mathbf{w}) = |\mathbf{w} \cdot \phi(x) - y|$$

- \mathbf{w} that minimizes training loss is **median** y
- **Median**: more robust to outliers