



Tutorial on Pytorch

Part 1

Review: CNN architecture

Classic CNN architecture:



Roadmap

- PyTorch basics
- PyTorch CNN

Install PyTorch in CPU mode

- 1) Install Anaconda (python environment manager)
- 2) Browse pytorch.org
- 3) Select and get installation command:

PyTorch Build	Stable (1.11.0)	Preview (Nightly)	LTS (1.8.2)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python	C++ / Java		
Compute Platform	CUDA 10.2	CUDA 11.3	ROCm 4.2 (beta)	CPU
Run this Command:	conda install pytorch torchvision torchaudio cpuonly -c pytorch			

- 4) Create a python environment with Anaconda

```
conda create -n your_env_name python=x.x
```

- 5) Activate the environment `Windows: activate your_env_name`
- 6) Install PyTorch with the command

Install PyTorch in CPU mode

- *Alternatively you can install PyTorch offline
- *You can install in GPU mode for fast training
- Recommended IDE:
 - PyCharm community (free) 
 - VSCode (free) 

Check if installation is successful:

```
>>> import torch
>>> import torchvision
>>> print(torch.__version__)
1.10.2
```

What's PyTorch

- It's a **Python-based** scientific computing package targeted at two sets of audiences:
 - A replacement for NumPy to use the power of **GPUs**
 - a **deep learning research platform** that provides maximum flexibility and speed
 - PyTorch CPU is slow, but it's easier to install and good for learning (**Coding is conceptually the same as GPU**)

Tensor

- Tensors are similar to NumPy's ndarray
- Tensors support fast computation using GPU

Construct a 5x3 matrix:

```
import torch
x = torch.empty(5, 3)
print(x)
```

Out:

```
>>> x = torch.empty(5,3)
>>> print(x)
tensor([[0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.]])
```

Tensor

Construct a randomly initialized matrix:

```
x = torch.rand(5, 3)  
print(x)
```

Out:

```
tensor([[0.6519, 0.6639, 0.5846],  
        [0.5429, 0.5386, 0.6401],  
        [0.5687, 0.1522, 0.1158],  
        [0.8838, 0.4172, 0.3364],  
        [0.4867, 0.5550, 0.8775]])
```

Construct a matrix filled zeros and of dtype long:

```
x = torch.zeros(5, 3, dtype=torch.long)  
print(x)
```

Out:

```
tensor([[0, 0, 0],  
        [0, 0, 0],  
        [0, 0, 0],  
        [0, 0, 0],  
        [0, 0, 0]])
```

Construct a tensor directly from data:

```
x = torch.tensor([5.5, 3])  
print(x)
```

Out:

```
tensor([5.5000, 3.0000])
```


Tensor

Create a tensor based on an existing tensor. These methods will **reuse properties** of the input tensor, e.g. dtype, unless new values are provided by user

```
x = x.new_ones(5, 3, dtype=torch.double)    # new_* methods take in sizes
print(x)
```

```
x = torch.randn_like(x, dtype=torch.float)  # override dtype!
print(x)                                    # result has the same size
```

Out:

```
tensor([[1., 1., 1.],
        [1., 1., 1.],
        [1., 1., 1.],
        [1., 1., 1.]], dtype=torch.float64)
tensor([[ 0.5863, -1.1109,  0.5939],
        [-0.4763, -1.5667, -1.1475],
        [ 0.0997,  0.8229,  0.2515],
        [ 3.2231, -1.0759,  0.0389],
        [ 0.5416, -0.6303,  0.9335]])
```

Get its size: `print(x.size())`

Out:

```
torch.Size([5, 3])
```

Operations

There are multiple syntaxes for operations.

Addition: syntax 1

```
y = torch.rand(5, 3)
print(x + y)
```

```
tensor([[ 0.7257, -0.6739,  1.4463],
        [ 0.2272, -0.8705, -0.2774],
        [ 0.4370,  1.7662,  0.5940],
        [ 3.3525, -0.6376,  0.5068],
        [ 0.6681, -0.0581,  1.7418]])
```

Addition: providing an output tensor as argument

```
result = torch.empty(5, 3)
torch.add(x, y, out=result)
print(result)
```

```
tensor([[ 0.7257, -0.6739,  1.4463],
        [ 0.2272, -0.8705, -0.2774],
        [ 0.4370,  1.7662,  0.5940],
        [ 3.3525, -0.6376,  0.5068],
        [ 0.6681, -0.0581,  1.7418]])
```

Addition: syntax 2

```
print(torch.add(x, y))
```

```
tensor([[ 0.7257, -0.6739,  1.4463],
        [ 0.2272, -0.8705, -0.2774],
        [ 0.4370,  1.7662,  0.5940],
        [ 3.3525, -0.6376,  0.5068],
        [ 0.6681, -0.0581,  1.7418]])
```

Addition: **in-place**

```
y.add_(x) # adds x to y
print(y)
```

```
tensor([[ 0.7257, -0.6739,  1.4463],
        [ 0.2272, -0.8705, -0.2774],
        [ 0.4370,  1.7662,  0.5940],
        [ 3.3525, -0.6376,  0.5068],
        [ 0.6681, -0.0581,  1.7418]])
```

Operations

- Any operation that mutates a tensor **in-place** is post-fixed with an **_**. For example: `x.copy_(y)`, `x.t_()`, **will change x**.
- You can use standard NumPy-like indexing with all bells and whistles!

```
print(x[:, 1])
```

Output: `tensor([-1.1109, -1.5667, 0.8229, -1.0759, -0.6303])`

- Resizing: If you want to resize/reshape tensor, you can use **`torch.view`**

```
x = torch.randn(4, 4)
y = x.view(16)
z = x.view(-1, 8) # the size -1 is inferred from other dimensions
print(x.size(), y.size(), z.size())
```

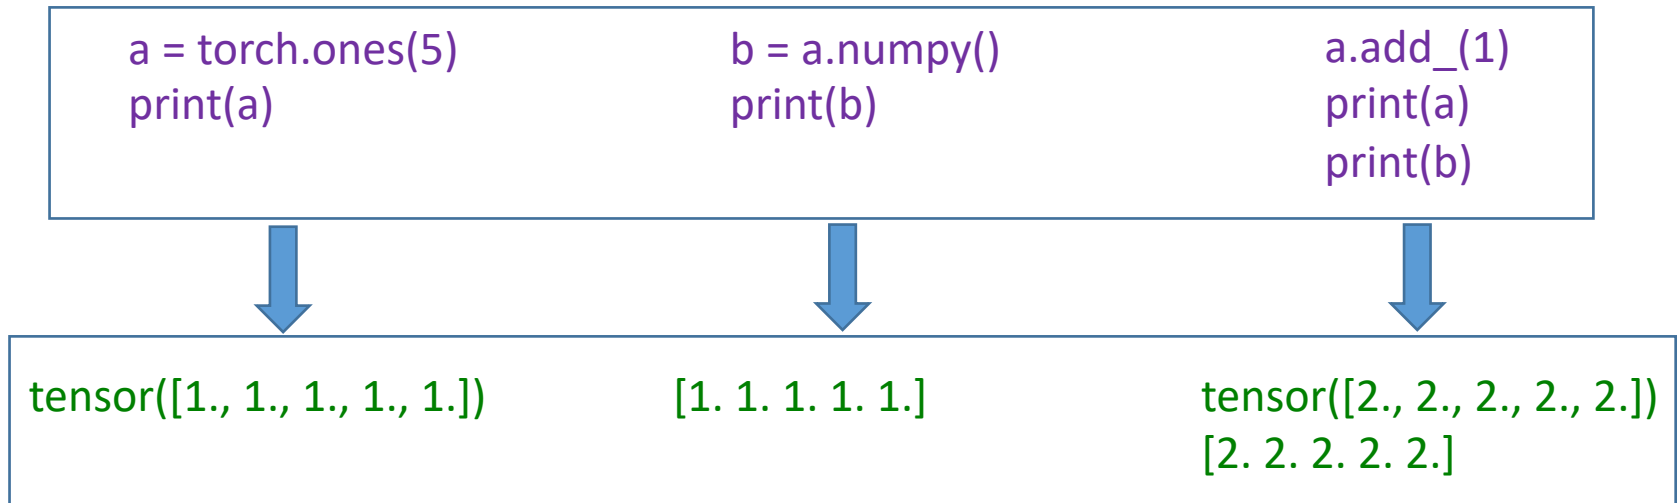
Output: `torch.Size([4, 4]) torch.Size([16]) torch.Size([2, 8])`

PyTorch supports 100+ Tensor operations, including transposing, indexing, slicing, mathematical operations, linear algebra, random numbers, etc., are described [here](#).

NumPy bridge

Converting a Torch Tensor to a NumPy array and vice versa is a breeze. The Torch Tensor and NumPy array will **share their underlying memory** locations (if the Torch Tensor is on **CPU**), and changing one will change the other.

Converting a Torch Tensor to a NumPy Array:



Converting NumPy array to Torch tensor

See how changing the np array changed the Torch Tensor automatically:

```
import numpy as np

a = np.ones(5)
b = torch.from_numpy(a)
np.add(a, 1, out=a)
print(a)
print(b)
```

Out:

```
[2. 2. 2. 2. 2.]
tensor([2., 2., 2., 2., 2.], dtype=torch.float64)
```

CUDA tensors

Tensors can be moved onto any device using the **.to** method.

```
# let us run this cell only if CUDA is available
# We will use ``torch.device`` objects to move tensors in and out of GPU
x = torch.randn(1)
print(x)
if torch.cuda.is_available():
    device = torch.device("cuda")      # a CUDA device object
    y = torch.ones_like(x, device=device) # directly create a tensor on GPU
    x = x.to(device)                   # or just use strings ``.to("cuda")``
    z = x + y
    print(z)
    print(z.to("cpu", torch.double))   # ``.to`` can also change dtype together!
```



Out:

```
tensor([0.8843])
tensor([1.8843], device='cuda:0')
tensor([1.8843], dtype=torch.float64)
```

PyTorch auto-gradient

Recall the computation graph we used:

- Tensors in PyTorch can be viewed as a **variables**
- PyTorch tracks all operations on tensors if you set **.requires_grad** as **True**
- After finishing all operations, just call **.backward** to compute all gradients automatically



Auto-gradient

Create a tensor and set `requires_grad=True` to track computation with it:

```
x = torch.ones(2, 2, requires_grad=True)
print(x)
```

Out:

```
tensor([[1., 1.],
        [1., 1.]], requires_grad=True)
```


Do a tensor operation:

```
y = x + 2
print(y)
```

y was created as a result of an operation, so it has a `grad_fn`

Out:

```
tensor([[3., 3.],
        [3., 3.]], grad_fn=<AddBackward0>)
```



Auto-gradient

Do more operations on y:

```
z = y * y * 3  
out = z.mean()  
print(z, out)
```

Out:

```
tensor([[27., 27.],  
        [27., 27.]], grad_fn=<MulBackward0>)  
tensor(27., grad_fn=<MeanBackward0>)
```

`.requires_grad_(...)` changes an existing Tensor's `requires_grad` flag in place.
The input flag defaults to **False** if not given.

```
a = torch.randn(2, 2)  
a = ((a * 3) / (a - 1))  
print(a.requires_grad)  
a.requires_grad_(True)  
print(a.requires_grad)  
b = (a * a).sum()  
print(b.grad_fn)
```

Out:

```
False  
True  
<SumBackward0 object at 0x7fcb3e741b00>
```

Back propagation with auto-gradient

Let's backpropagate now by calling `out.backward()`

```
out.backward()
```

Print gradients $d(\text{out})/dx$

```
print(x.grad)
```

Out:


```
tensor([[4.5000, 4.5000],  
        [4.5000, 4.5000]])
```

How to verify the output? $o = \frac{1}{4} \sum_i z_i \quad z_i = 3(x_i + 2)^2 \quad z_i|_{x_i=1} = 27$

$$\frac{\partial o}{\partial x_i} = \frac{3}{2}(x_i + 2) \quad \frac{\partial o}{\partial x_i}|_{x_i=1} = \frac{9}{2} = 4.5$$

Demo

Recall the neural network we used:



Intermediate hidden units:


$$h_j = \text{ReLU}(\mathbf{v}_j \cdot \phi(x)) \quad \text{ReLU}(z) = \max(z, 0)$$

Output:

$$\text{score}(\phi(x); \mathbf{v}, \mathbf{w}) = \mathbf{w} \cdot [h_1, h_2]$$

How to compute implement this network with PyTorch?

A mathematical way



$$\frac{\partial}{\partial w_1} \text{score}(\phi(x); \mathbf{v}, \mathbf{w}) = \frac{\partial \sum_{j=1}^2 w_j \text{ReLU}(\mathbf{v}_j \cdot \phi(x))}{\partial w_1} = \text{ReLU}(\mathbf{v}_1 \cdot \phi(x))$$

$$\left. \begin{array}{l} \phi(x) = [1.0, 0.5, 0.2] \\ \mathbf{v} = \left[\underbrace{[0.2, 0.0, 0.5]}_{\mathbf{v}_1}, \underbrace{[0.0, 1.0, 0.5]}_{\mathbf{v}_2} \right] \\ \mathbf{w} = [0.5, 0.5] \end{array} \right\} \longrightarrow \frac{\partial}{\partial w_1} \text{score}(\phi(x); \mathbf{v}, \mathbf{w}) = 0.3$$

PyTorch implementation

```
C:\Users\30535\AppData\Local\Programs  
score: 0.45  
gradient w.r.t. w: 0.30, 0.60  
tensor([0.3000, 0.6000])  
gradient w.r.t. v:  
tensor([[0.5000, 0.2500, 0.1000],  
        [0.5000, 0.2500, 0.1000]])
```

[live solution]

Summary

- PyTorch is able to track all computations automatically associated with tensors (variables)
- PyTorch is able to calculate all gradients automatically

Acknowledgement

Reference and thanks to:

- **Sandford University CS221 Course:**

Artificial Intelligence: Principles and Techniques

<https://stanford-cs221.github.io/autumn2022/>

- **PyTorch Introduction**

<https://courses.cs.washington.edu/courses/cse446/19au/section9.html>