



# Tutorial on Pytorch

Part 2

# Review: CNN architecture

## Classic CNN architecture:





# Review: training procedure of CNN

- Define network architecture
- Iteratively get a batch of training data
  - Forward propagation to compute loss (error)
  - Backward propagation to compute gradient
  - Update weights:  $\text{weight} = \text{weight} - \text{learning\_rate} * \text{gradient}$

# Review: PyTorch auto-gradient

Recall the computation graph we used:


- Tensors in PyTorch can be viewed as a **variables**
- PyTorch tracks all operations on tensors if you set **.requires\_grad** as **True**
- After finishing all operations, just call **.backward** to compute all gradients automatically



# Roadmap

- PyTorch basics
- PyTorch CNN

# Define the network



## Required modules:

```
import torch
```

```
import torch.nn as nn
```

```
import torch.nn.functional as F
```

Depends on autograd to define models and differentiate them

Contains useful functions such as ReLU

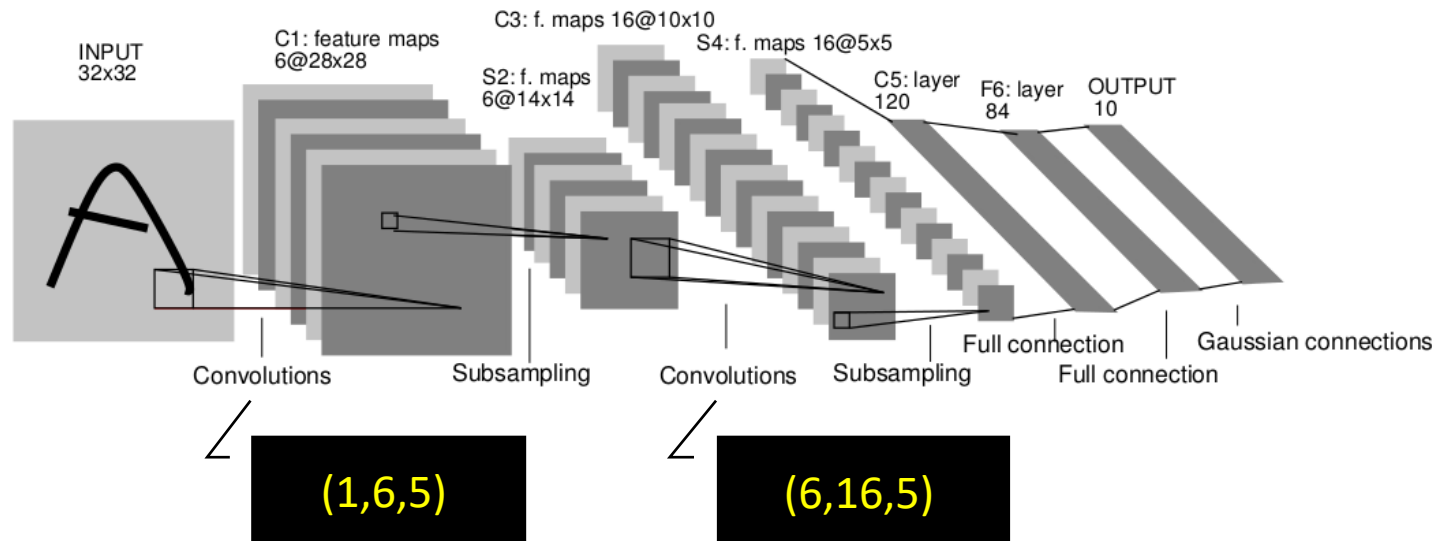
# Define the net: convolution **object**

```
nn.Conv2d(in_channels, out_channels, kernel_size,  
          stride=1, padding=0, bias=True)
```

- in\_channels: the **number of channels** of the input feature map
- out\_channels: the **number of channels** of the output feature map
- kernel\_size: (h, w), actually is **h x w x in\_channels**
- stride: (s\_h, s\_w)
- padding: (p\_h, p\_w)
- bias: with (**True**) or without (**False**) bias

$$w_{out} = \lfloor (w_{in} - k + 2 \times p) / s + 1 \rfloor$$


$$h_{out} = \lfloor (h_{in} - k + 2 \times p) / s + 1 \rfloor$$



# Define the network: activation and pooling functions

`F.relu(input, inplace=False)`

`F.max_pool2d(input, kernel_size, stride=None, padding=0)`




`F.max_pool2d(6@28x28, (2,2))`

`F.max_pool2d(16@10x10, (2,2))`



# Define the network: FC layer **object**

```
nn.Linear(in_feature, out_features, bias=True)
```



# Demo

[LeNet demo]



Net Class

Output:

```
Net
(  
  (conv1):Conv2d(1, 6, kernel_size=(3, 3), stride=(1, 1))  
  (conv2): Conv2d(6, 16, kernel_size=(3, 3), stride=(1, 1))  
  (fc1): Linear(in_features=400, out_features=120, bias=True)  
  (fc2): Linear(in_features=120, out_features=84, bias=True)  
  (fc3): Linear(in_features=84, out_features=10, bias=True)  
)
```

# Forward propagation

```
net = Net()
input = torch.randn(1, 1, 32, 32)
out = net(input)
print(out)
```

- torch.nn only support inputs that are a mini-batch of samples, in a form of `nSamples x nChannels x Height x Width` (e.g., `1 x 1 x 32 x 32`)
- Net object is callable with a `__call__` method inherited from `torch.nn.modules.module.Module`, which performs **forward passing**

Output:

```
tensor([[ 0.0484,  0.1541, -0.0510,  0.1475,  0.0433,  0.1539, -0.0636, -0.0635,
          0.0050, -0.0705]], grad_fn=<AddmmBackward>)
```

# Loss function

Recall: loss function takes the (output, target) pair of inputs, and computes a value that estimates **how far away the output is from the target**.

There are a bunch of loss functions to use in PyTorch, for example here:

```
net = Net()
input = torch.randn(1, 1, 32, 32)
out = net(input)
print(out)
```

```
target = torch.randn(10) # a dummy target, for example
target = target.view(1, -1) # make it the same shape as output
criterion = nn.MSELoss()
loss = criterion(output, target)
print(loss)
```

```
criterion = nn.MSELoss()
loss = criterion(output, target)
print(loss)
```

# Track the computation graph

Now, if you follow loss in the backward direction, using its `.grad_fn` attribute, you will see a graph of computations that looks like:

```
input -> conv2d -> relu -> maxpool2d -> conv2d -> relu  
      -> maxpool2d -> linear -> relu -> linear -> relu -> linear  
      -> log_softmax -> nll_loss
```

So, when we call `loss.backward()`, the whole graph is differentiated w.r.t. the loss, and all Tensors in the graph that has `requires_grad=True` will have their `.grad` tensor **accumulated** with the gradient.



Remember we are processing a batch of examples

# Backward propagation

To backpropagate the error all we have to do is to implement `loss.backward()`. You need to **clear the existing gradients** though, **else gradients will be accumulated to existing gradients**.

```
net.zero_grad()
print('conv1.bias.grad before backward')
print(net.conv1.bias.grad)
loss.backward()
print('conv1.bias.grad after backward')
print(net.conv1.bias.grad)
```

Output:

```
conv1.bias.grad before backward
None
conv1.bias.grad after backward
tensor([-0.0198, 0.0211, -0.0077, -0.0060, 0.0203, -0.0385])
```

# Update the weights

The simplest update rule used in practice is the Stochastic Gradient Descent (SGD):

$$\text{weight} = \text{weight} - \text{learning\_rate} \times \text{gradient}$$

We can implement this using simple Python code:

```
learning_rate = 0.01
for f in net.parameters():
    f.data.sub_(f.grad.data * learning_rate)
```

# Update the weights

To enable other optimizers, you need to import `torch.optim`:

```
import torch.optim as optim
```

```
# create your optimizer
```

```
optimizer = optim.SGD(net.parameters(), lr=0.01)
```

```
# in your training loop:
```

```
optimizer.zero_grad() # zero the gradient buffers
```

```
output = net(input)
```

```
loss = criterion(output, target)
```

```
loss.backward()
```

```
optimizer.step() # Does the update
```



# Hand-written digits recognition



[demo]

# Summary

- Define PyTorch network by inheriting the `nn.Module` Class
- PyTorch processes data in mini-batch
- Training routine in PyTorch: forward propagation, backward propagation, weights updating
- CNN achieves the state-of-the-art performance on image classification